Assignment 2
Markus Gall
Carmina Coronel

How does SGD work?
Gradient Descent is used to minimize loss. The goal is to find the point where the loss function has its global minimum (ideally = 0). Stochastic Gradient Descent (SGD) is an iterative approach to do this by processing the training data in mini batches of size $S = 1$. The algorithm for SGD is the Minibatch Gradient Descent, where a subset of the training data, sampled randomly, is processed to calculate the model error and update the model coefficients. In practice, using $S > 1$ is common and still referred to as SGD. In deep learning, common S values are 64, 128, 256.

What are your findings on the example data (part1)?
Madsen was the easiest example to find the global minima since only no other minimas are present. With nearly each setting the algorithm find the minima faster or slower depending on learning rate and the usage of nesterov momentum.
This was different for the Beale example. Even though different starting points found a minima, the starting point and the learning rate had influence on which one endpoint was found.
Schaffern4 is meant to confuse the algorithm and the algorithm is most likely to find a local minima only since many of these were present. Using a high learning rate may leads to a jump out of the local minima but finds itself most likely in another area and completely changes direction again.
In general, we find that using a high learning rate is faster but can make too many big "jumps" and may miss the minimas. On the other hand a with small learning rate, the algorithm proceeds very slowly. The starting point can also influence what minimum is reached. Using the nesterov momentum the algorithm processes faster because it builds up momentum when successive gradients are similar.

Which network architecture did you choose for part 2 and why?
For part 2, we use the following model architecture:
CONV1(activation=relu) → POOL → CONV1(activation=relu) → POOL → FLATTEN → FCL(activation=relu) → FCL(activation=relu) → SOFTMAX

*FCL - fully connected layer
Learning rate = 0.01
Momentum = 0.9

Results (epoch=100):
   train loss: 0.000 +- 0.000
   val acc: 0.726
   Best Accuracy of 0.7300342969132778 at epoch 74

We used CNN because for image classifications this type of network works well. It works well because it extracts features from the images using spatial relationship of the pixels. We initially chose this architecture as we felt a 2 level convolutional layer would give good results.

We then added another convolutional layer in an attempt to improve the results.

CONV1(activation=relu) → POOL → CONV1(activation=relu) → POOL → CONV1(activation=relu) → POOL → FLATTEN → FCL(activation=relu) → FCL(activation=relu) → SOFTMAX

Results (epoch=100):
        train loss: 0.000 +- 0.000
        val acc: 0.736
        Best Accuracy of 0.7491425771680549 at epoch 39

However results did not substantially improve.

Did you have problems reaching a low training error?
Training loss is almost 0 but our validation accuracy is around 75% suggesting possible overfitting and model can still be improved.

What are the goals of data augmentation, regularization, and early stopping?
The goals of data augmentation, regularization, and early stopping is to avoid underfitting and overfitting.

How exactly did you use these techniques (hyperparameters, combinations) and what were your results (train and val performance)?
We performed data augmentation (flipping with cropping, and blurring), with or without dropout layers in different locations, different weight decays, different location of dropout layers. We also played with different learning rates and batch normalization.

We improved our validation accuracy by using data augmentation (flipping, cropping of images) from 73% to 83%.

List all experiments and results, even if they did not work well, and discuss them.
        (results shown are at epoch 100)
- Data Augmentation (flipping and cropping)
    - Flipping and Cropping
      train loss: 0.297 +- 0.051
      val acc: 0.826
      Best Accuracy of 0.8280254777070064 at epoch 92
      Better validation accuracy by using flipping and cropping of images. If more data augmentation techniques is employed, results can still be improved.

    - Gaussing filters for blurring
      We tried adding Gaussian filters for blurring but it did not improve the validation accuracy.
- Weight Decay 0.001 and 0.0001

- (2) For 0.0001
    train loss: 0.309 +- 0.053
    val acc: 0.825
    Best Accuracy of 0.8358647721705047 at epoch 93

    Minimal change to results when weight decay is lowered. We can get away with 0.001 (default we used)
- Dropout Layers
    - Before the 2nd convolutional layer
        train loss: 0.402 +- 0.054
        val acc: 0.805
        Best Accuracy of 0.8054875061244487 at epoch 100
    - As first layer
        Acc: 0.78

    Dropout layers did not work for us. We tried putting it in different places in the model but no improvement, worst, it reduces our model results.
- Learning Rates (0.01, 0.001, 0.001) and batch normalization
    We also tried different learning rates (0.01, 0.001, 0.0001) and batch normalization after every convolution layer. We did not get better accuracy numbers.

If you utilized transfer learning, explain what you did and your results.
We utilized transfer learning by using densenet121. We trained the whole training set with no weights freeze. This is different from what was usually done, which is to freeze weights of the convolutional parts of the pretrained model while the fully connected layer is trained anew.

Using the pretrained model, we achieved a 95% validation and test accuracy. This model is now the model saved in best_model.